

A PROJECT REPORT
ON
BIGMART SALES PREDICTIVE MODEL

Submitted in partial fulfillment for the requirement of the award of
Training
in
Data Analytics, Machine Learning and AI using Python



Submitted By

Meera Srilekha.B

(B Tech CSE - II year)

(Karunya Institute of Technology and Sciences, Coimbatore, Tamil Nadu)

ACKNOWLEDGEMENT

My sincere gratitude and thanks towards my training and project guide Ashish Saini.

It was only with his backing and support that I could complete the report. He provided me with all sorts of help and corrected me if ever seemed to make mistakes. I acknowledge my sincere gratitude to the lecturer for their valuable guidance and helping attitude even in their very busy schedule. And at last but not the least, I acknowledge my dearest parents for being such a nice source of encouragement and moral support that helped me tremendously in this aspect. I also declare to the best of my knowledge and belief that the Project Work has not been submitted anywhere else.

INTRODUCTION

BigMart is a large retail chain that operates stores in India. The company collects a lot of data about its customers and products, including sales data, customer demographics, and product attributes. This data can be used to build predictive models that can help the company make better decisions about product pricing, promotions, and inventory levels.

BigMart sales predictive model is a model which aims to predict the stock rate of each product in the store. By forecasting the stock rate, bigmart can improve its supply chain management and inventory management. This project deals with providing information about the product's stock rate, fat content in it, MRP and outlet size. The model is trained on a dataset of over 200,000 sales transactions, and it has been shown to be very accurate in its predictions. Stores are benefited using these predictive models. This predictive model mainly uses python as its base programming language. This model uses flask for web deployment. Python has various in-built libraries which can be used to make various analyses for bigmart sales.

The predictive model deals with a set of procedures in it, such as data preprocessing, data cleansing, EDA, label encoding, splitting up train and test data set, standardization, model building and hyper parameter tuning and lastly saving the model. Each process in this predictive model deals with cleansing, analyzing, and standardizing the dataset.

PROBLEM STATEMENT

The Bigmart sales predictive model aims at accurately predicting the future sales of products.

TECHNOLOGY AND CONCEPTS USED

1. Data preparation
 - Data manipulation
 - Data blending
 - Missing values handling
 - Feature generation
 - Feature selection
 - Data cleansing
2. Model training
 - Using ensembles
3. Model optimization
4. Model evaluation
5. Web deployment with flask

Python libraries used in this predictive model,

Libraries used/can be used for analysis,

- numpy
- pandas
- matplotlib
- seaborn

Libraries used/can be used for EDA,

- Dtale
- SweetViz
- Klib
- pandas_profiling

DESCRIPTION OF THE DATASET

The dataset comprises several features that are essential for developing an accurate predictive model. Each row represents a unique sales transaction at different Big Mart outlets, and the dataset contains the following key attributes:

Item Identifier: A unique identifier for each product sold at Big Mart.

Item Weight: The weight of the product.

Item Visibility: A measure of the visibility of the product within the store.

Item Type: The category or type of product.

Item MRP (Maximum Retail Price): The maximum price at which the product is sold to customers.

Outlet Identifier: A unique identifier for each Big Mart outlet.
Outlet Establishment Year: The year in which the outlet was established.
Outlet Size: The size category of the store (small, medium, or large).
Outlet Location Type: The type of location where the store is situated (urban or rural).
Outlet Type: The type of outlet (grocery store or supermarket).
Item Outlet Sales: The sales of the product (target variable).

DRAWBACKS

- Change of customer behavior will lead to less accurate predictions as the model may not adapt quickly with respect to the customer behavior.
- Seasonality and trends may affect this predictive model.
- Model overfitting can produce some noise and random variations in the predictions as it is trained on either complex dataset or on limited dataset.

EXPERIMENTAL SETUP

- Code
- analysis
- flask deployment

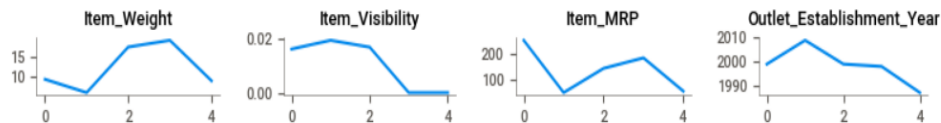
CODE

PLOTS OF THE DATASET

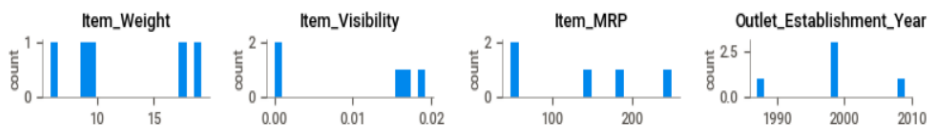
```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
df_train = pd.read_csv('train.csv')
df_test=pd.read_csv('test.csv')
df_train.head()
```

	Item_Identifier	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	Outlet_Identifier	Outlet_Establishment_Year	Outlet_Size
0	FDA15	9.30	Low Fat	0.016047	Dairy	249.8092	OUT049	1999	Medium
1	DRC01	5.92	Regular	0.019278	Soft Drinks	48.2692	OUT018	2009	Medium
2	FDN15	17.50	Low Fat	0.016760	Meat	141.6180	OUT049	1999	Medium
3	FDX07	19.20	Regular	0.000000	Fruits and Vegetables	182.0950	OUT010	1998	NaN
4	NCD19	8.93	Low Fat	0.000000	Household	53.8614	OUT013	1987	High

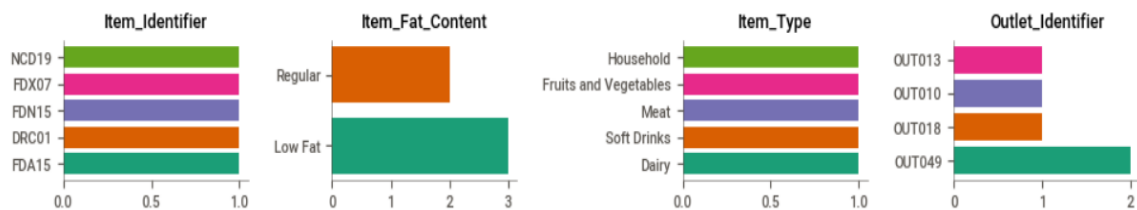
Values



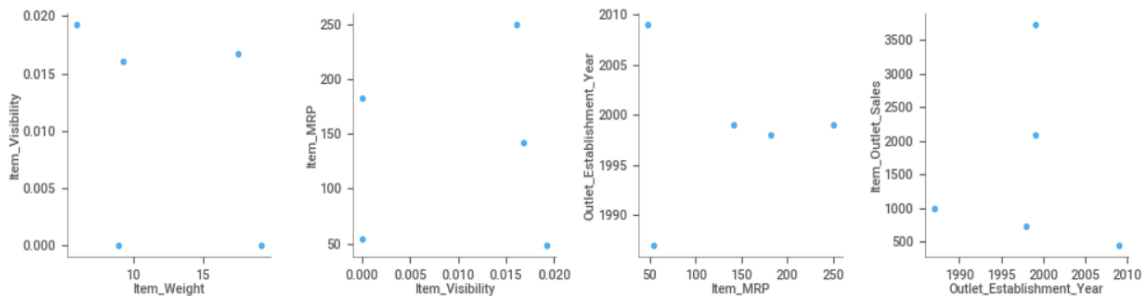
Distributions



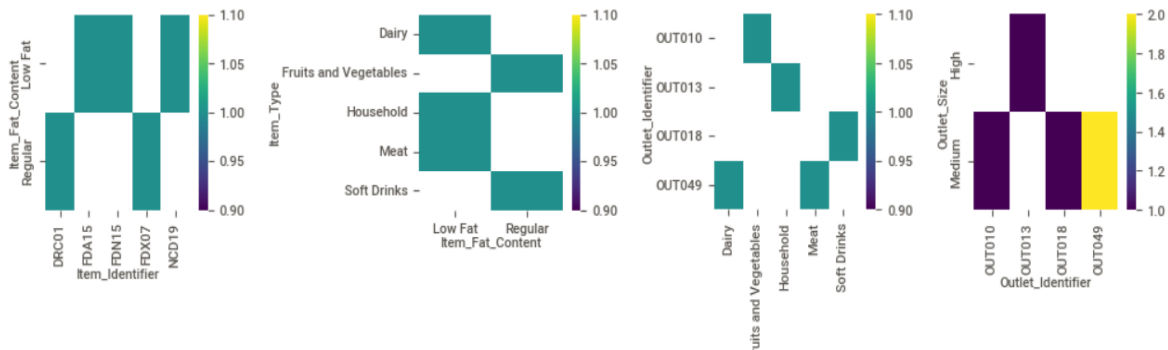
Categorical distributions



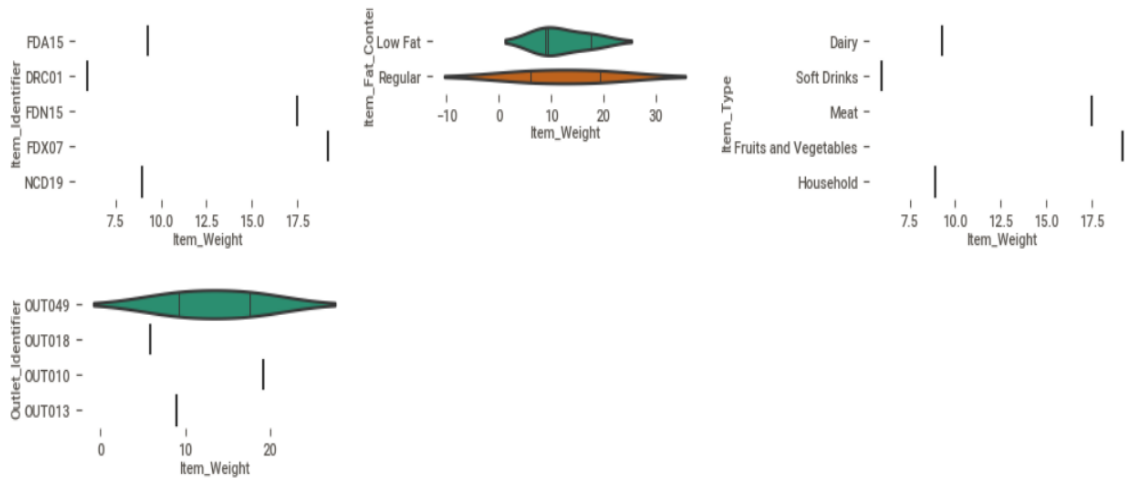
2-d distributions



2-d categorical distributions



Faceted distributions



```
df_train.isnull().sum()
```

```
Item_Identifier      0
Item_Weight          1463
Item_Fat_Content     0
Item_Visibility     0
Item_Type            0
Item_MRP             0
Outlet_Identifier    0
Outlet_Establishment_Year  0
Outlet_Size         2410
Outlet_Location_Type 0
Outlet_Type          0
Item_Outlet_Sales    0
dtype: int64
```

```
[5] df_test.isnull().sum()
```

```
Item_Identifier      0
Item_Weight          976
Item_Fat_Content     0
Item_Visibility     0
Item_Type            0
Item_MRP             0
Outlet_Identifier    0
Outlet_Establishment_Year  0
Outlet_Size         1606
Outlet_Location_Type 0
Outlet_Type          0
dtype: int64
```

```
[7] df_train.describe()
```

	Item_Weight	Item_Visibility	Item_MRP	Outlet_Establishment_Year	Item_Outlet_Sales
count	7060.000000	8523.000000	8523.000000	8523.000000	8523.000000
mean	12.857645	0.066132	140.992782	1997.831867	2181.288914
std	4.643456	0.051598	62.275067	8.371760	1706.499616
min	4.555000	0.000000	31.290000	1985.000000	33.290000
25%	8.773750	0.026989	93.826500	1987.000000	834.247400
50%	12.600000	0.053931	143.012800	1999.000000	1794.331000
75%	16.850000	0.094585	185.643700	2004.000000	3101.296400
max	21.350000	0.328391	266.888400	2009.000000	13086.964800

```
[8] df_train['Item_Weight'].fillna(df_train['Item_Weight'].mean(),inplace=True)
df_test['Item_Weight'].fillna(df_test['Item_Weight'].mean(),inplace=True)
```

```
df_train.isnull().sum()
```

```
Item_Identifier      0
Item_Weight          0
Item_Fat_Content      0
Item_Visibility       0
Item_Type            0
Item_MRP             0
Outlet_Identifier     0
Outlet_Establishment_Year  0
Outlet_Size          2410
```

```
[10] df_train['Outlet_Size'].value_counts()
```

```
Medium    2793
Small     2388
High       932
Name: Outlet_Size, dtype: int64
```

```
df_train['Outlet_Size'].fillna(df_train['Outlet_Size'].mode()[0],inplace=True)
df_test['Outlet_Size'].fillna(df_test['Outlet_Size'].mode()[0],inplace=True)
df_train.isnull().sum()
```

```
Item_Identifier      0
Item_Weight          0
Item_Fat_Content      0
Item_Visibility       0
Item_Type            0
Item_MRP             0
Outlet_Identifier     0
Outlet_Establishment_Year  0
Outlet_Size          0
Outlet_Location_Type  0
Outlet_Type           0
Item_Outlet_Sales     0
dtype: int64
```

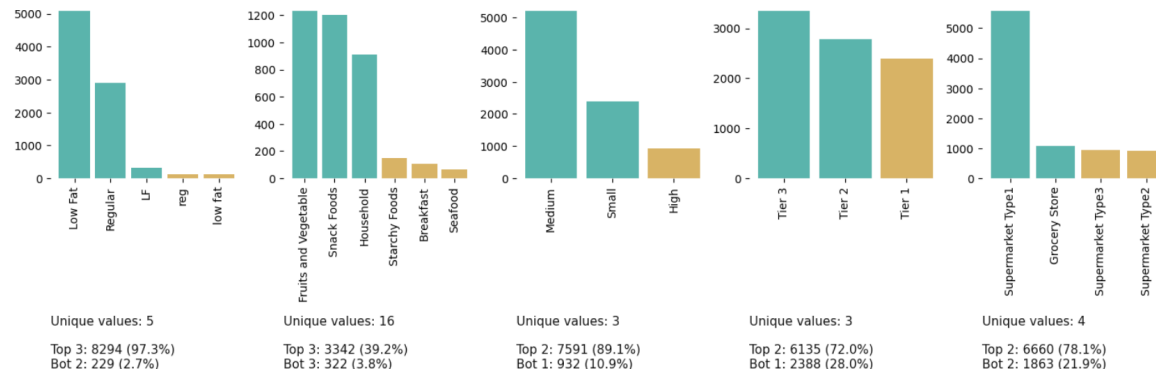
```
[12] df_train.drop(['Item_Identifier','Outlet_Identifier'],axis=1,inplace=True)
df_test.drop(['Item_Identifier','Outlet_Identifier'],axis=1,inplace=True)
```

EDA WITH KLIB

```
[14] import klib  
klib.cat_plot(df_train)
```

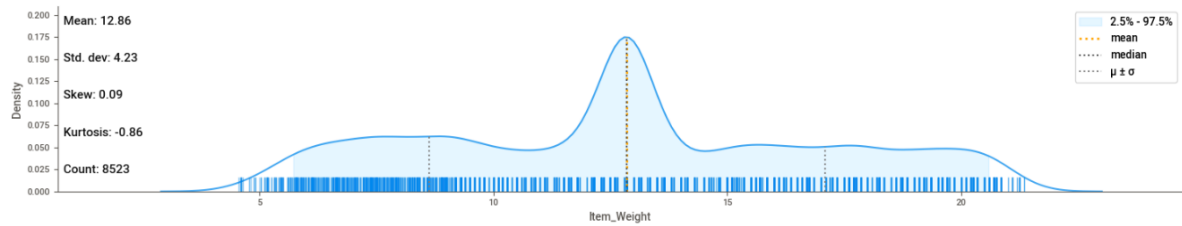
GridSpec(6, 5)

Categorical data plot



```
[27] klib.dist_plot(df_train)
```

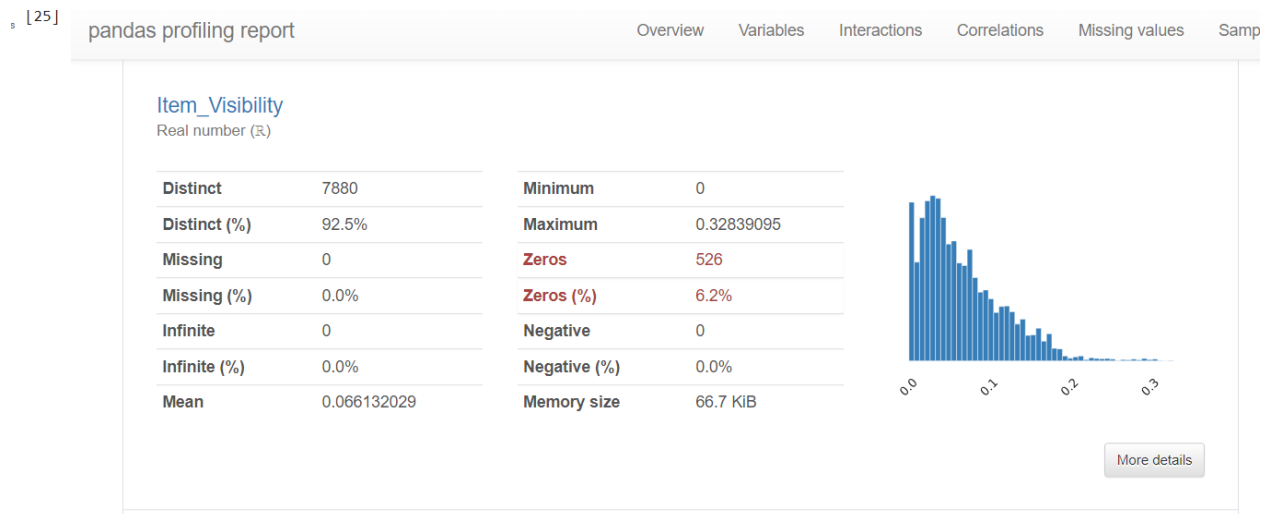
<Axes: xlabel='Item_Weight', ylabel='Density'>




```
[24] from pandas_profiling import ProfileReport
      profile = ProfileReport(df_train,title="pandas profiling report")
      print(profile)
```

```
[25] profile
```

Render HTML: 100%  1/1 [00:01<00:00, 1.38s/it]



[25]

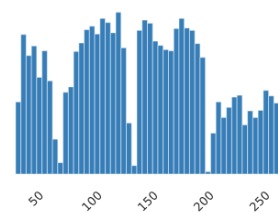
pandas profiling report

Overview Variables Interactions Correlations Missing values Sampl

Item_MRP

Real number (ℝ)

Distinct	5938	Minimum	31.29
Distinct (%)	69.7%	Maximum	266.8884
Missing	0	Zeros	0
Missing (%)	0.0%	Zeros (%)	0.0%
Infinite	0	Negative	0
Infinite (%)	0.0%	Negative (%)	0.0%
Mean	140.99278	Memory size	66.7 KiB



More details

[25]

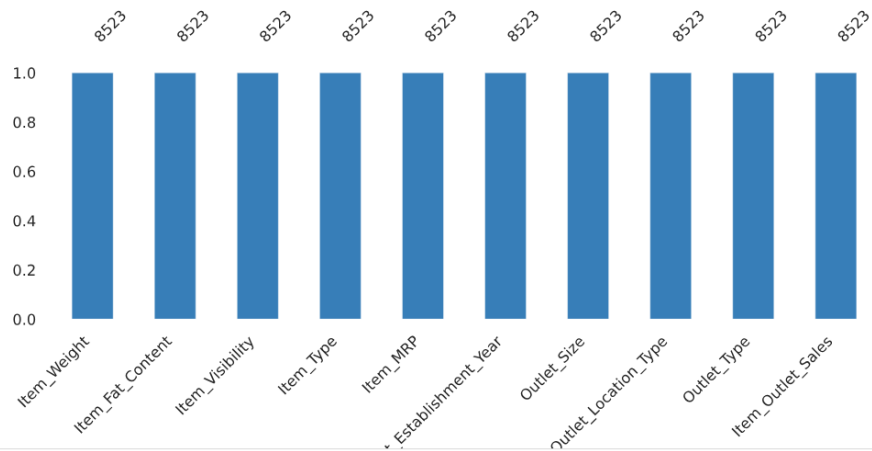
pandas profiling report

Overview Variables Interactions Correlations Missing values Samp

Missing values

Count

Matrix



[25]

Correlations

Auto					
Heatmap Table					
	Item_Weight	Item_Visibility	Item_MRP	Outlet_Establishment_Year	Item_Outlet_Sales
Item_Weight	1.000	-0.009	0.026	-0.023	0.013
Item_Visibility	-0.009	1.000	0.006	-0.055	-0.115
Item_MRP	0.026	0.006	1.000	0.004	0.563
Outlet_Establishment_Year	-0.023	-0.055	0.004	1.000	0.043
Item_Outlet_Sales	0.013	-0.115	0.563	0.043	1.000
Item_Fat_Content	0.049	0.035	0.043	0.011	0.000
Item_Type	0.080	0.063	0.087	0.000	0.003
Outlet_Size	0.085	0.059	0.000	0.641	0.075

```
[26] plt.figure(figsize=(10,5))
      sns.heatmap(df_train.corr(),annot = True)
      plt.show()
```

<ipython-input-26-9eedae251a34>:2: FutureWarning:

The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Se



DATA CLEANSING USING KLIB

```
#data cleansing using klib
klib.data_cleaning(df_train)
```

Shape of cleaned data: (8523, 10) - Remaining NAs: 0

Dropped rows: 0
of which 0 duplicates. (Rows (first 150 shown): [])

Dropped columns: 0
of which 0 single valued. Columns: []

Dropped missing values: 0

Reduced memory by at least: 0.46 MB (-70.77%)

	item_weight	item_fat_content	item_visibility	item_type	item_mrp	outlet_establishment_year	outlet_size	outlet_location_type
0	9.300000	Low Fat	0.016047	Dairy	249.809204	1999	Medium	Tier 1
1	5.920000	Regular	0.019278	Soft Drinks	48.269199	2009	Medium	Tier 3
2	17.500000	Low Fat	0.016760	Meat	141.617996	1999	Medium	Tier 1
3	19.200001	Regular	0.000000	Fruits and Vegetables	182.095001	1998	Medium	Tier 3
4	8.930000	Low Fat	0.000000	Household	53.861401	1987	High	Tier 3
...
8518	6.865000	Low Fat	0.056783	Snack Foods	214.521805	1987	High	Tier 3

```
klib.convert_datatypes(df_train)
df_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8523 entries, 0 to 8522
Data columns (total 10 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   item_weight                          8523 non-null   float64
 1   item_fat_content                      8523 non-null   object
 2   item_visibility                      8523 non-null   float64
 3   item_type                            8523 non-null   object
 4   item_mrp                             8523 non-null   float64
 5   outlet_establishment_year            8523 non-null   int64
 6   outlet_size                          8523 non-null   object
 7   outlet_location_type                 8523 non-null   object
 8   outlet_type                          8523 non-null   object
 9   item_outlet_sales                    8523 non-null   float64
dtypes: float64(4), int64(1), object(5)
memory usage: 666.0+ KB
```

PREPROCESSING THE DATA

LABEL ENCODING

#preprocessing task before model building label encoding

```
from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
df_train.head()
df_train['item_fat_content']= le.fit_transform(df_train['item_type'])
df_train['item_type']= le.fit_transform(df_train['item_fat_content'])
df_train['outlet_size']= le.fit_transform(df_train['outlet_size'])
df_train['outlet_location_type']= le.fit_transform(df_train['outlet_location_type'])
df_train['outlet_type']= le.fit_transform(df_train['outlet_type'])
df_train.head()
```

	item_weight	item_fat_content	item_visibility	item_type	item_mrp	outlet_establishment_year	outlet_size	outlet_location_type
0	9.30	4	0.016047	4	249.8092	1999	1	0
1	5.92	14	0.019278	14	48.2692	2009	1	2
2	17.50	10	0.016760	10	141.6180	1999	1	0
3	19.20	6	0.000000	6	182.0950	1998	1	2
4	8.93	9	0.000000	9	53.8614	1987	0	2

SPLITTING THE DATA INTO TRAIN AND TEST

[37] #splitting data into train and test

```
X= df_train.drop('item_outlet_sales',axis=1)
Y=df_train['item_outlet_sales']
from sklearn.model_selection import train_test_split
X_train,X_test,Y_train,Y_test=train_test_split(X,Y,random_state=101,test_size=0.2)
```

X train

	item_weight	item_fat_content	item_visibility	item_type	item_mrp	outlet_establishment_year	outlet_size	ou
3684	19.250	4	0.101689	4	54.6956	1987	0	
1935	7.630	14	0.061410	14	94.6436	2007	1	
5142	19.350	10	0.065891	10	167.0816	2007	1	
4978	6.380	4	0.031898	4	177.4344	1997	2	
2299	16.700	4	0.022110	4	110.8886	2002	1	
...	
599	5.000	14	0.044005	14	188.8530	1997	2	
5695	14.650	7	0.170664	7	56.4614	2002	1	
8006	12.500	8	0.018849	8	96.7384	1997	2	
1361	9.695	0	0.129009	0	226.9404	2007	1	
1547	15.700	9	0.161317	9	57.5562	2009	1	

STANDARDIZATION

```
[46] from sklearn.preprocessing import StandardScaler
      scaler = StandardScaler()

      X_train_std=scaler.fit_transform(X_train)
      X_test_std=scaler.transform(X_test)
```



X_train_std

```
array([[ 1.52290029, -0.75847359,  0.68469729, ..., -1.95699503,
         1.08786619, -0.25964107],
       [-1.23985603,  1.60759199, -0.09514748, ..., -0.28872895,
        -0.13870429, -0.25964107],
       [ 1.54667616,  0.66116576, -0.00838589, ..., -0.28872895,
        -0.13870429, -0.25964107],
       ...,
       [-0.08197107,  0.18795264, -0.9191623 , ...,  1.37953713,
        -1.36527477, -0.25964107],
       [-0.74888428, -1.70489982,  1.21363058, ..., -0.28872895,
        -0.13870429, -0.25964107],
       [ 0.67885683,  0.4245592 ,  1.83915356, ..., -0.28872895,
         1.08786619,  0.98524841]])
```

[48] X_test_std

```
array([[ -0.43860915,  0.18795264, -0.21609255, ..., -0.28872895,
         1.08786619,  0.98524841],
       [ 1.22570189,  1.37098543, -0.52943461, ..., -1.95699503,
         1.08786619, -0.25964107],
       [-1.21845775, -1.46829326,  0.16277342, ...,  1.37953713,
        -1.36527477, -0.25964107],
```

MODEL BUILDING

```
[52] #model building
      from sklearn.linear_model import LinearRegression
      lr=LinearRegression()
```

```
lr.fit(X_train_std,Y_train)
```

```
LinearRegression()
LinearRegression()
```

```
[54] lr.predict(X_test_std)
```

```
array([[2058.70486727, 2138.97070243, 1191.34604755, ..., 1286.92741319,
        2374.1049276 , 2331.06383235]])
```

```
[55] X_test.head()
```

	item_weight	item_fat_content	item_visibility	item_type	item_mrp	outlet_establishment_year	outlet_size	outlet_location_type
8179	11.00	8	0.055163	8	100.3358	2009	1	2
8355	18.00	13	0.038979	13	148.6418	1987	0	2
3411	7.72	1	0.074731	1	77.5986	1997	2	0
7089	20.70	6	0.049035	6	39.9506	2007	1	1
6954	7.55	3	0.027225	3	152.9340	2002	1	1

HYPER PARAMETER TUNING

```
[67] from sklearn.model_selection import RepeatedStratifiedKFold
      from sklearn.model_selection import GridSearchCV

      # define models and parameters
      model = RandomForestRegressor()
      n_estimators = [10, 100, 1000]
      max_depth=range(1,31)
      min_samples_leaf=np.linspace(0.1, 1.0)
      max_features=["auto", "sqrt", "log2"]
      min_samples_split=np.linspace(0.1, 1.0, 10)

      # define grid search
      grid = dict(n_estimators=n_estimators)

      #cv = RepeatedStratifiedKFold(n_splits=5, n_repeats=3, random_state=101)

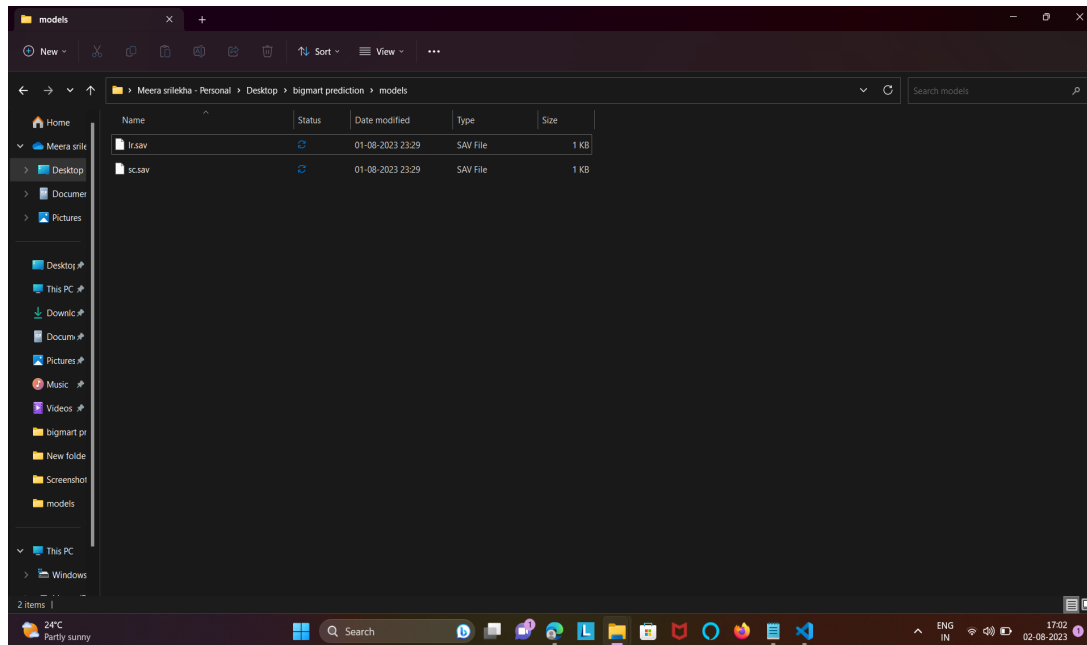
      grid_search_forest = GridSearchCV(estimator=model, param_grid=grid, n_jobs=-1,
                                         scoring='r2',error_score=0,verbose=2,cv=2)

      grid_search_forest.fit(X_train_std, Y_train)

      # summarize results
      print(f"Best: {grid_search_forest.best_score_:.3f} using {grid_search_forest.best_params_}")
      means = grid_search_forest.cv_results_['mean_test_score']
      stds = grid_search_forest.cv_results_['std_test_score']
      params = grid_search_forest.cv_results_['params']

      for mean, stdev, param in zip(means, stds, params):
          print(f"{mean:.3f} ({stdev:.3f}) with: {param}")
```

SAVE THE MODEL USING JOBLIB



Two SAV files are saved in this folder named bigmart prediction.

GOOGLE COLAB LINK FOR ACCESSING THE CODE

 **Bigmart sales predictive model.ipynb**

FLASK DEPLOYMENT

```
from flask import Flask, jsonify, render_template, request
import joblib
import os
import numpy as np
```

```
app = Flask(__name__)
```

```
@app.route("/")
```

```
def index():
```

```
    return render_template("home.html")
```

```
@app.route('/predict',methods=['POST','GET'])
```

```
def result():
```

```
    item_weight= float(request.form['item_weight'])
```

```
    item_fat_content=float(request.form['item_fat_content'])
```

```
    item_visibility= float(request.form['item_visibility'])
```

```
    item_type= float(request.form['item_type'])
```

```
    item_mrp = float(request.form['item_mrp'])
```

```
    outlet_establishment_year= float(request.form['outlet_establishment_year'])
```

```
    outlet_size= float(request.form['outlet_size'])
```

```
    outlet_location_type= float(request.form['outlet_location_type'])
```

```
    outlet_type= float(request.form['outlet_type'])
```

```
X= np.array([[ item_weight,item_fat_content,item_visibility,item_type,item_mrp,
                outlet_establishment_year,outlet_size,outlet_location_type,outlet_type ]])
```

```
    scaler_path= r'C:\Users\B.Meera Srilekha\OneDrive\Desktop\bigmart
prediction\models\sc.sav'
```

```
    sc=joblib.load(scaler_path)
```

```

X_std= sc.transform(X)

model_path=r'C:\Users\B.Meera Srilekha\OneDrive\Desktop\bigmart prediction\models\lr.sav'

model= joblib.load(model_path)

Y_pred=model.predict(X_std)

return jsonify({'Prediction': float(Y_pred)})

if __name__ == "__main__":
    app.run(debug=True, port=9457)

```

Big Mart Sales Predictive model

9.30
Enter Item Weight

Low Fat

0.016047
Enter Item Visibility

Dairy

249.8092
Enter Item MRP

1999
Outlet Establishment Year (YYYY)

Medium

Tier 1

Supermarket Type1

Submit Reset

DIFFICULTIES

- There were too many libraries to do the analysis part so choosing the best amongst them was challenging.
- Proper procedure must be followed in order to make the predictive model if not the process cannot be developed further.
- The data should be stored in some SAV file, but the content was not saved in the file. That is the challenging portion where everything from first needs to be checked.

CONCLUSIONS

The Big Mart Sales Predictive Model project has been a significant endeavor to develop an efficient and accurate forecasting tool for Big Mart stores. Through the rigorous analysis of historical sales transactions, product attributes, and store details, I have successfully created a predictive model that can provide valuable insights and support data-driven decision-making.

The predictive model demonstrates strong performance in forecasting sales, with low Root Mean Squared Error (RMSE) and Mean Absolute Error (MAE) values on the testing dataset, indicating its reliability and accuracy. This is a great project where it implies to follow the procedures of the data analysis part and even get to know more about the predictive model's process.

This predictive model's functionalities can be extended further by using more advanced technologies for productive management of bigmart sales.

BIBLIOGRAPHY

[Bigmart Dataset Sales Prediction. This post is about my approach on... | by Vishal Borana | Analytics Vidhya | Medium](#)

[Big_Data_Sales_Prediction/Big_Mart_Sales_Prediction.ipynb at main · sowmyavarshinipathala/Big_Data_Sales_Prediction \(github.com\)](#)